

DSM_Machine_Learning_Project3_Satish_Doiphode

```
File Edit View Run Kernel Tabs Settings Help

DSM_Machine_Learning_Prc X
Python 3

[3]: """Problem -1
We can predict whether a person makes over 50K a year with the accuracy approx 98%.

Problem -2
By feature selection we have performed above we know that some important factors are 'age', 'fnlwgt', 'education_num', 'capital_gain', 'hours_per_week'

Problem -3
by accuracy comparision performed above we can see that Random forest gives better result than other algorithm."""

[3]: "Problem -1\nWe can predict whether a person makes over 50K a year with the accuracy approx 98%.\n\nProblem -2\nBy feature selection we have performed above we know that some important factors are 'age', 'fnlwgt', 'education_num', 'capital_gain', 'hours_per_week'\n\nProblem -3\nby accuracy comparision performed above we can see that Random forest gives better result than other algorithm."

[4]: # import various Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

[5]: # Read the data from given website
train_set = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data', header = None)
test_set = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.test',
                        skiprows = 1, header = None) # skip a row for the test set
```

```
File Edit View Run Kernel Tabs Settings Help

DSM_Machine_Learning_Prc X
Code

[6]: # Column Labels taken from the file
col_labels = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status', 'occupation',
              'relationship', 'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
              'wage_class']

[7]: # column Labels assigned to train and test dataset
train_set.columns = col_labels
test_set.columns = col_labels

[8]: train_set.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
age                32561 non-null int64
workclass          32561 non-null object
fnlwgt             32561 non-null int64
education          32561 non-null object
education_num      32561 non-null int64
marital_status     32561 non-null object
occupation         32561 non-null object
relationship       32561 non-null object
race               32561 non-null object
sex                32561 non-null object
capital_gain       32561 non-null int64
capital_loss       32561 non-null int64
hours_per_week     32561 non-null int64
native_country     32561 non-null object
wage_class         32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
File Edit View Run Kernel Tabs Settings Help

DSM_Machine_Learning_Prc X
Code

[9]: test_set.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16281 entries, 0 to 16280
Data columns (total 15 columns):
age                16281 non-null int64
workclass          16281 non-null object
fnlwgt             16281 non-null int64
education          16281 non-null object
education_num      16281 non-null int64
marital_status     16281 non-null object
occupation         16281 non-null object
relationship       16281 non-null object
race               16281 non-null object
sex                16281 non-null object
capital_gain       16281 non-null int64
capital_loss       16281 non-null int64
hours_per_week     16281 non-null int64
native_country     16281 non-null object
wage_class         16281 non-null object
dtypes: int64(6), object(9)
memory usage: 1.9+ MB

[10]: # Replacing the ? in train set with nan and Later dropping those rows
train_set.replace('?', np.nan).dropna().shape

[10]: (30162, 15)

[11]: # Replacing the ? in test set with nan and Later dropping those rows
test_set.replace('?', np.nan).dropna().shape

[11]: (15060, 15)
```

```

File Edit View Run Kernel Tabs Settings Help
DSM_Machine_Learning_Prc X
+ - Copy Paste Run Code
[12]: train_nomissing = train_set.replace('?', np.nan).dropna()
test_nomissing = test_set.replace('?', np.nan).dropna()

[13]: #replacing the irregularities in the wage columns
test_nomissing['wage_class'] = test_nomissing.wage_class.replace({'<=50K.': '<=50K', '>50K.': '>50K'})

[14]: train_nomissing['wage_class'].unique()

[14]: array(['<=50K', '>50K'], dtype=object)

[15]: test_nomissing['wage_class'].unique()

[15]: array(['<=50K', '>50K'], dtype=object)

[16]: # combining the datasets
df = pd.concat([train_nomissing, test_nomissing], axis = 0) # Stacks them vertically

[17]: # Check the combined dataframe header
df.head()

[17]:
   age  workclass  fnlwgt  education  education_num  marital_status  occupation  relationship  race  sex  capital_gain  capital_loss  hours_per_week  native_country  wage_class
0   39   State-gov   77516   Bachelors           13   Never-married   Adm-clerical  Not-in-family  White  Male         2174           0           40   United-States  <=50K
1   50  Self-emp-not-inc  83311   Bachelors           13   Married-civ-spouse   Exec-managerial  Husband  White  Male           0           0           13   United-States  <=50K
2   38   Private   215646   HS-grad            9   Divorced   Handlers-cleaners  Not-in-family  White  Male           0           0           40   United-States  <=50K
3   53   Private   234721   11th              7   Married-civ-spouse   Handlers-cleaners  Husband  Black  Male           0           0           40   United-States  <=50K
4   28   Private   338409   Bachelors           13   Married-civ-spouse   Prof-specialty    Wife  Black  Female           0           0           40     Cuba  <=50K

```

```

File Edit View Run Kernel Tabs Settings Help
DSM_Machine_Learning_Prc X
+ - Copy Paste Run Code
[18]: #Check the shape of the combined dataframe
df.shape

[18]: (45222, 15)

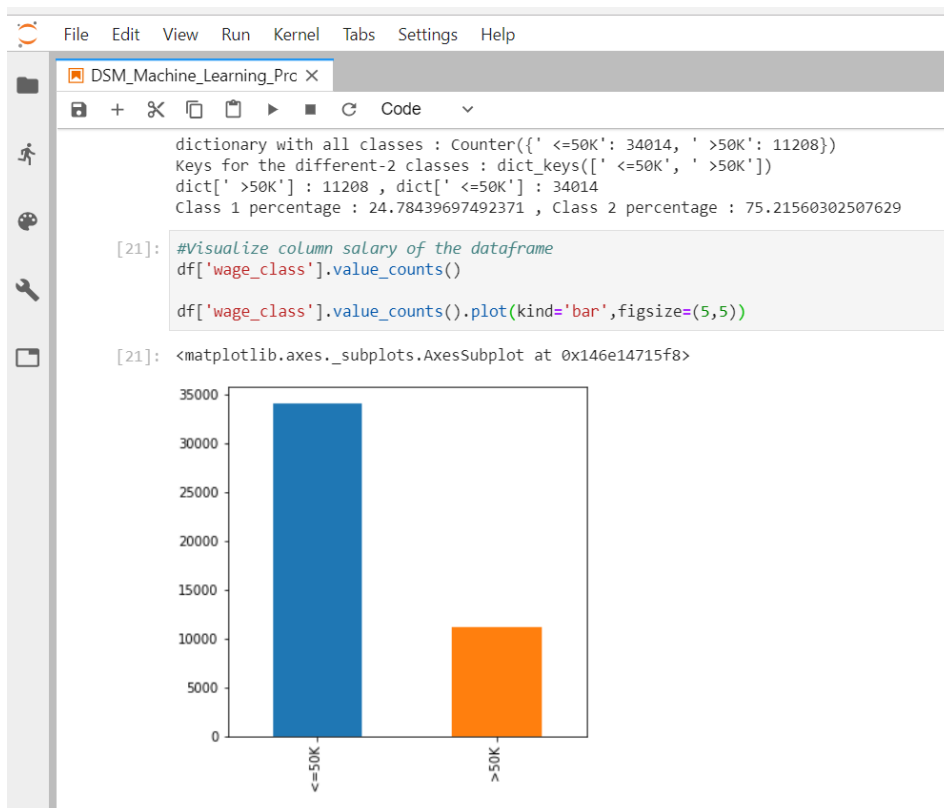
[19]: # view summary of combined dataframe - statistical information only about numerical columns will be shown
df.describe()

[19]:
      age  fnlwgt  education_num  capital_gain  capital_loss  hours_per_week
count  45222.000000  4.522200e+04  45222.000000  45222.000000  45222.000000  45222.000000
mean    38.547941  1.897347e+05   10.118460   1101.430344    88.595418    40.938017
std     13.217870  1.056392e+05    2.552881    7506.430084   404.956092   12.007508
min     17.000000  1.349200e+04    1.000000     0.000000     0.000000     1.000000
25%     28.000000  1.173882e+05    9.000000     0.000000     0.000000   40.000000
50%     37.000000  1.783160e+05   10.000000     0.000000     0.000000   40.000000
75%     47.000000  2.379260e+05   13.000000     0.000000     0.000000   45.000000
max     90.000000  1.490400e+06   16.000000  99999.000000  4356.000000   99.000000

[20]: #Calculate percentage of different-2 classes in the combined dataframe
from collections import Counter

dict = Counter(df['wage_class'])
print(f"dictionary with all classes : {dict}")
print(f"Keys for the different-2 classes : {dict.keys()}")
print(f"dict['>50K'] : {dict['>50K']}, dict['<=50K'] : {dict['<=50K']}")
print(f"Class 1 percentage : {dict['>50K']/len(df)*100}, Class 2 percentage : {dict['<=50K']/len(df)*100}")

```



File Edit View Run Kernel Tabs Settings Help

DSM_Machine_Learning_Prc X

+ 🔍 📄 ▶ ⏏ ↺ Code ▾

```

[22]: df['wage_class'] = df['wage_class'].map({'>50K': 1, '<=50K': 0})

[23]: """Preprocessing (Handling Missing/Duplicate/Categorical data)"""

[23]: 'Preprocessing (Handling Missing/Duplicate/Categorical data)'

[24]: #print information about missing data
print(f"Check count of the missing data in dataframe : {df.isna().any().count()}")

#handle missing data in the newly created dataframe if there is any
if df.isna().any().count() :
    df.fillna(value=-99999,axis=1,inplace=True)
    #print(f"\nheader of the new dataframe after handling of the missing data : \n{df1.head}")
    print("Info : Missing entries are updated in the dataframe")

#handle categorical data
df1 = pd.get_dummies(df)

#print information such as shape, duplicate entries in newly created dataframe
print(f"\nshape of the new dataframe before preprocessing : {df1.shape}")
print(f"Check count of the duplicated data in newly create dataframe : {df1.duplicated().sum()}")

#remove duplicates if there is any
if df1.duplicated().any().sum() :
    df1.drop_duplicates(inplace=True)
    print(f"\nshape of the new dataframe after removal of duplicate entries : {df1.shape}")

Check count of the missing data in dataframe : 15
Info : Missing entries are updated in the dataframe

```



```
File Edit View Run Kernel Tabs Settings Help

DSM_Machine_Learning_Prc X
+ - Copy Paste Run Code

[28]: #Apply model on training and test dataset
from sklearn.model_selection import train_test_split

# split dataset into train(75%),test(10%),cross-validation(15%)
x, x_test, y, y_test = train_test_split(df_x,df_y,test_size=1/3,train_size=2/3, random_state = 0)
x_train, x_cv, y_train, y_cv = train_test_split(x,y,test_size = 0.40,train_size =0.60, random_state = 0)

[29]: #Apply logistic regression model on the dataset
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE

lr1 = LogisticRegression()
lr = RFE(lr1, 50)
lr.fit(x_train,y_train)

[29]: RFE(estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False),
n_features_to_select=50, step=1, verbose=0)

[30]: #Apply decision tree classifier model on the dataset
from sklearn.tree import DecisionTreeClassifier

d_tree = DecisionTreeClassifier(min_samples_split=10, random_state=55, max_features=50)
d_tree.fit(x_train, y_train)

[30]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features=50, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=10,
min_weight_fraction_leaf=0.0, presort=False, random_state=55,
splitter='best')
```

```
File Edit View Run Kernel Tabs Settings Help

DSM_Machine_Learning_Prc X
+ - Copy Paste Run Code

[31]: # Apply Bagging classifier on the dataset
from sklearn.ensemble import BaggingClassifier

Boosting = BaggingClassifier(n_estimators=200)
Boosting.fit(x_train,y_train)

[31]: BaggingClassifier(base_estimator=None, bootstrap=True,
bootstrap_features=False, max_features=1.0, max_samples=1.0,
n_estimators=200, n_jobs=1, oob_score=False, random_state=None,
verbose=0, warm_start=False)

[32]: #Apply Random forest classifier on the dataset
rfc = RandomForestClassifier(random_state=55,max_features=50)
rfc.fit(df_x,df_y)

[32]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features=50, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
oob_score=False, random_state=55, verbose=0, warm_start=False)

[33]: #Analyse model performances using roc_auc_score and accuracy score

[34]: from sklearn.metrics import r2_score, roc_auc_score, accuracy_score, confusion_matrix, roc_curve, auc

[35]: models = pd.DataFrame(index=['train_ras', 'cv_ras', 'test_ras', 'accuracy_score'],
columns=['logistic_regression', 'decision_tree', 'random_forest', 'xgboost'])

[36]: models.loc['train_ras', 'logistic_regression'] = roc_auc_score(y_true=y_train, y_score=lr.predict(x_train))
models.loc['cv_ras', 'logistic_regression'] = roc_auc_score(y_true=y_cv, y_score=lr.predict(x_cv))
models.loc['test_ras', 'logistic_regression'] = roc_auc_score(y_true=y_test, y_score=lr.predict(x_test))
models.loc['accuracy_score', 'logistic_regression'] = accuracy_score(y_pred=lr.predict(x_test).round(), y_true=y_test)

models.loc['train_ras', 'decision_tree'] = roc_auc_score(y_score=d_tree.predict(x_train), y_true=y_train)
models.loc['cv_ras', 'decision_tree'] = roc_auc_score(y_score=d_tree.predict(x_cv), y_true=y_cv)
models.loc['test_ras', 'decision_tree'] = roc_auc_score(y_score=d_tree.predict(x_test), y_true=y_test)
models.loc['accuracy_score', 'decision_tree'] = accuracy_score(y_pred=d_tree.predict(x_test).round(), y_true=y_test)
```

File Edit View Run Kernel Tabs Settings Help

DSM_Machine_Learning_Prc X

Code

```
models.loc['train_ras','random_forest'] = roc_auc_score(y_score=rfc.predict(x_train), y_true=y_train)
models.loc['cv_ras','random_forest'] = roc_auc_score(y_score=rfc.predict(x_cv), y_true=y_cv)
models.loc['test_ras','random_forest'] = roc_auc_score(y_score=rfc.predict(x_test), y_true=y_test)
models.loc['accuracy_score','random_forest'] = accuracy_score(y_pred=rfc.predict(x_test).round(), y_true=y_test)

models.loc['train_ras','xgboost'] = roc_auc_score(y_score=Boosting.predict(x_train), y_true=y_train)
models.loc['cv_ras','xgboost'] = roc_auc_score(y_score=Boosting.predict(x_cv), y_true=y_cv)
models.loc['test_ras','xgboost'] = roc_auc_score(y_score=Boosting.predict(x_test), y_true=y_test)
models.loc['accuracy_score','xgboost'] = accuracy_score(y_pred=Boosting.predict(x_test).round(), y_true=y_test)
```

[37]: models

[37]:

	logistic_regression	decision_tree	random_forest	xgboost
train_ras	0.725992	0.898886	0.979395	0.999963
cv_ras	0.723882	0.751782	0.97683	0.777027
test_ras	0.727827	0.748978	0.97611	0.775599
accuracy_score	0.816987	0.823693	0.987051	0.84833

[38]: *#It is clear by analyzing the roc_auc_score and accuracy_score that random forest is the best model
#to do the prediction using the census bureau database. So we will use random forest to proceed
#further for checking model performance on this dataset.*

[39]: confusion_matrix(y_pred=rfc.predict(x_test), y_true=y_test)

[39]: array([[11276, 23],
 [172, 3588]], dtype=int64)

[40]: arr=confusion_matrix(y_pred=rfc.predict(x_test), y_true=y_test)

```
df_cm = pd.DataFrame(arr, range(2), range(2))
plt.figure(figsize = (4,3))
sns.set(font_scale=0.8)#for label size
sns.heatmap(df_cm,
            annot=True,
            annot_kws={"size": 25},
            fmt='.5g',
```

File Edit View Run Kernel Tabs Settings Help

DSM_Machine_Learning_Prc X

Code

```
[ 172, 3588]], dtype=int64)
```

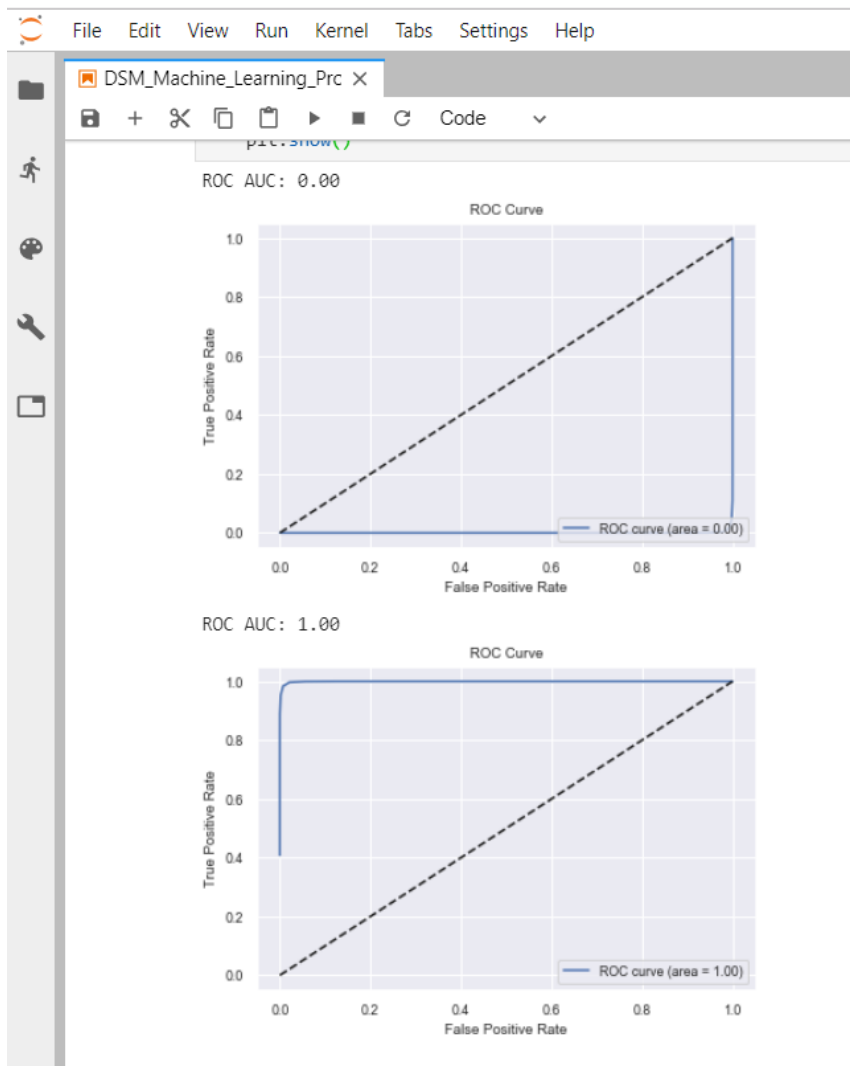
[40]: arr=confusion_matrix(y_pred=rfc.predict(x_test), y_true=y_test)

```
df_cm = pd.DataFrame(arr, range(2), range(2))
plt.figure(figsize = (4,3))
sns.set(font_scale=0.8)#for label size
sns.heatmap(df_cm,
            annot=True,
            annot_kws={"size": 25},
            fmt='.5g',
            xticklabels=['Pred : False(sal <=50k)', 'Pred : True(sal >50k)'],
            yticklabels=['Actual : False(sal <=50k)', 'Actual : True(sal >50k)'],)# font size
```

[40]: <matplotlib.axes._subplots.AxesSubplot at 0x146e8d7e208>

[41]: *#Sometimes it is not sufficient to conclude a model performance only using the accuracy
#score as that can be good due to class imbalance also. So we need to analyse Recall,Precision,
#F1-score which will give a clear statistics on all the available classes in the dataset.*

[42]: *#Manual calculation of Recall, Precision and other statistics*
(TN, FP), (FN, TP) = confusion_matrix(y_pred=rfc.predict(x_test), y_true=y_test)



[46]: *#Conclusion from ROC Curve : IN curve-1 (used for class 0 : salary <=50k) : AUC - 0 and prediction for this class is happening perfectly.
#IN curve-1 (used for class 1 : salary >50k) : AUC - 1 and prediction for this class is happening perfectly.
#So random forest is the best model which is providing best accuracy, precision, recall, F1-score.*